

Experiences with the super learner

Thomas Alexander Gerds

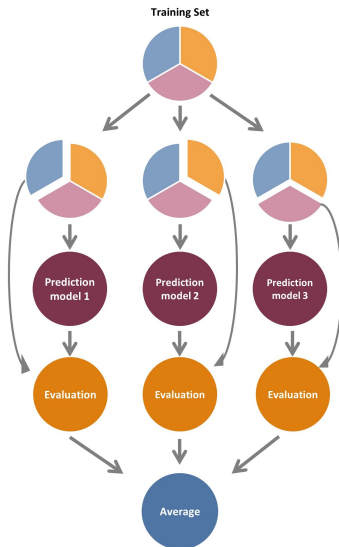
Cross-validation

Prediction performance describes how well a model will do on future patients.

Unfortunately, this is never knowable.

The best we can do is to simulate the model being applied to future patients by **repeatedly** splitting the data set into training and validation part.

3-fold cross-validation



Outline

Frequently asked questions about cross-validation:

- Q1: 5-fold, 10-fold, 20-fold or bootstrap?
- Q2: Bootstrap with or without replacement?
- Q4: Model hyper parameter selection using all data?
- Q5: In a survival setting where IPCW is used to estimate performance/loss: estimate censoring distribution in each training set separately?
- Q6: Use prediction performance to create a new learner?

Super learner

The super learner¹ (aka stacked regression²) uses cross-validation data (aka level-1 data) to combine multiple prediction models into a super model.

From that perspective, the super learner is just another machine learning algorithm which takes in the data and spits out a predicted risk.

The idea is to pre-specify all steps of modelling and to deal with misspecified regression models by controlling all data-dependent modeling steps with cross-validation.

¹van der Laan 2007

²Wolpert 1992, Breiman 1996

Two semi-parametric problems – in survival analysis

Prediction of an event between time 0 and time t based on predictors X :

$$P(T \leq t | X = x)$$

Average treatment effect:

$$\int \left\{ P(T \leq t | X = x, A = 1) - P(T \leq t | X = x, A = 0) \right\} P(dx)$$

Two semi-parametric problems – in survival analysis

Prediction of an event between time 0 and time t based on predictors X :

$$P(T \leq t | X = x)$$

Average treatment effect:

$$\int \left\{ P(T \leq t | X = x, A = 1) - P(T \leq t | X = x, A = 0) \right\} P(dx)$$

Both problems require a risk prediction model and hence a super learner could be applied . . .

Modeling strategy

Training data $D_n = \{\tilde{T}_i, \Delta_i, X_i\}_{i=1}^n$

Cox/logistic regression
Backward elimination
Lars and his three cousins
Cart and RandomForests
Support vector machines
Recursive Neural Networks
Super learner

$D_n \rightarrow$

\rightarrow

$\underbrace{R(D_n)}$

Risk prediction model

Risk prediction model

Black box

$$R(D_n)(t, X_{new}) \approx P(T \leq t | X_{new})$$

Cox/logistic regression
Backward elimination
Lars and his three cousins
Cart and RandomForests
Deep Learning
Recursive Neural Networks
regression

$$(t, X_{new}) \rightarrow \underbrace{R(D_n)(t, X_{new})}_{\text{predicted risk of event}}$$

Candidate modelling strategies

Flexible Cox regression with penalty:

$$R_n(t, X) = \exp(-\hat{\Lambda}_0(t) \exp(\hat{\beta}_1 \hat{f}_1(X_1) + \cdots + \hat{\beta}_p \hat{f}_p(X_p)))$$

where restricted cubic splines with k -knots are applied to the continuous predictor variables

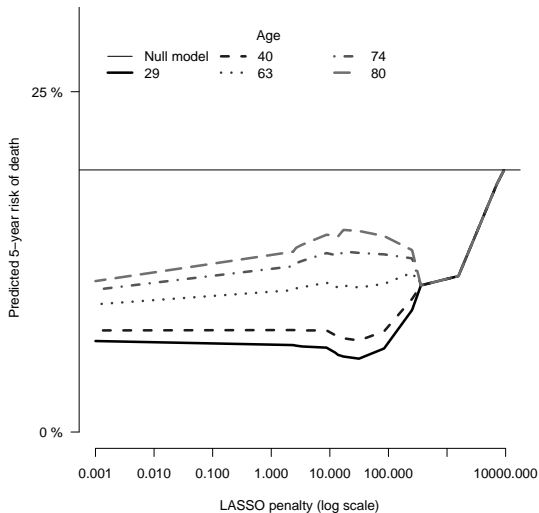
$$\hat{f}(x) = \hat{\alpha}_0 + \hat{\alpha}_1 x + \hat{\alpha}_2 \left\{ (x - t_j)_+^3 - (x - t_{k-1})^3 + (t_k - t_j)/(t_k - t_{k-1}) \right\}$$

where the partial likelihood estimate of the hazard ratios

$$\hat{\beta} = \hat{\beta}(\hat{\lambda})$$

depends on the shrinkage parameter λ .

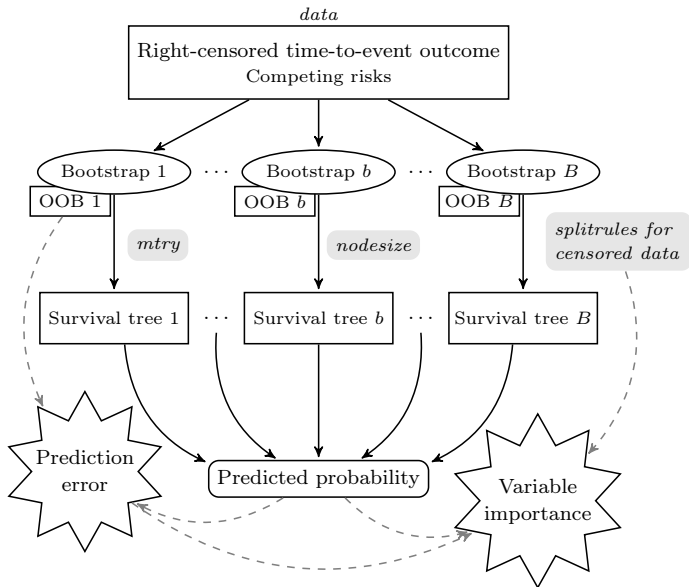
Penalized Cox regression




Q4: Model hyper parameter selection using all data?

A4: No, do not use all the data to optimize the penalty parameter $\hat{\lambda}$. Instead, use some form of cross-validation.

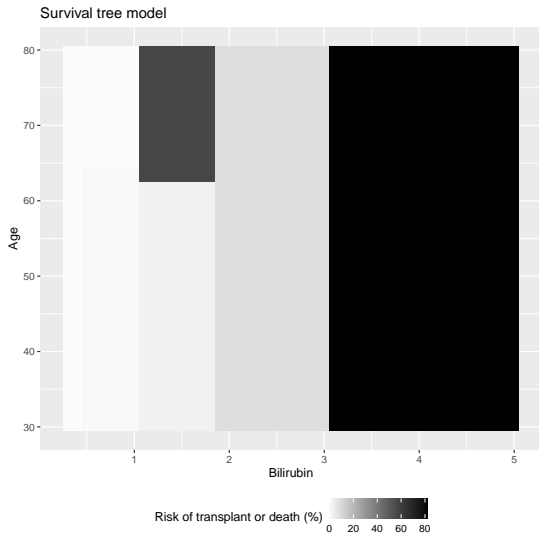
Candidate modelling strategies: random survival forest




Predicted 3-year risk: survival tree in bootstrap sample



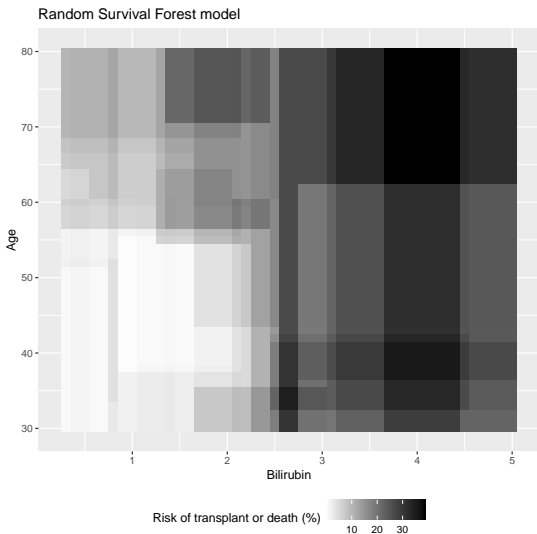
Predicted 3-year risk: survival tree in another bootstrap sample



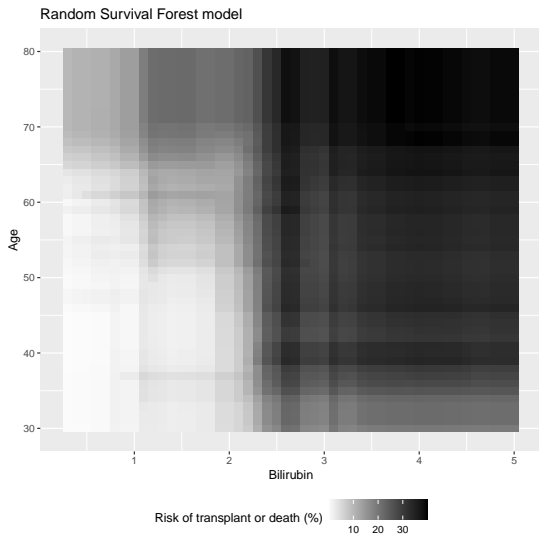
Predicted 3-year risk: survival tree in another bootstrap



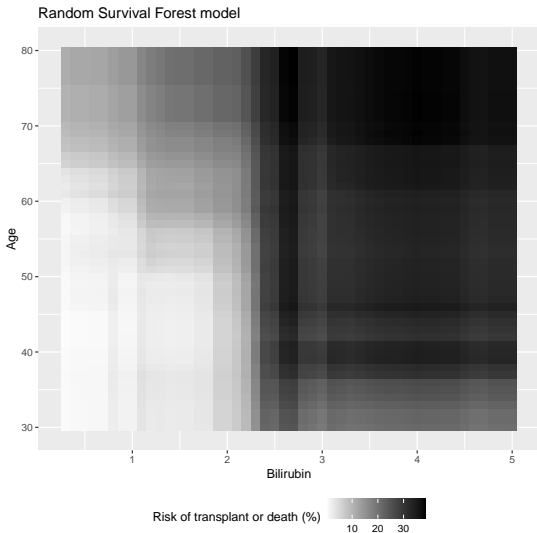
Predicted 3-year risk: random forest with 10 survival trees



Predicted 3-year risk: random forest with 100 survival trees



Predicted 3-year risk: random forest with 1000 survival trees



Q2: Bootstrap with or without replacement?

A2: Both candidate modelling strategies internally use cross-validation or bootstrap to build prediction models. Hence, any outer cross-validation procedure should use bootstrap without replacement.

Current situation

We have developed two prediction models (Cox and random survival forest) both estimate the risk of an event until a given time t in new patients.

Which one is better?

The aim is to assess the prediction performance of the model in new patients. But: there are **no data** of new patients.

Fundamental idea

Data splitting is very intuitive:

we hide one part of the data, learn on the rest, and then check our *knowledge* on what was hidden.

There is a hidden parameter here: how much we hide and how much we show.

Right censored data

Notation:

n	sample size
$m < n$	bootstrap subsample size
$\tilde{T} = \min(T, C)$	time to event
$\Delta = 1\{T \leq C\}$	event indicator
X	p-dimensional covariate
t	prediction horizon
$D_n = \{(\tilde{T}_i, \Delta_i, X_i)\}_{i=1}^n$	Dataset
$D_{b,m}^* = \{(\tilde{T}_{i,b}^*, \Delta_{i,b}^*, X_{i,b}^*)\}_{i=1}^m$	Bootstrap dataset $m < n$
$R_t(D_n)$	trained prediction model (size n)
$R_t(D_{b,m}^*)$	trained prediction model (size m)
$R_t(D_n)(X_{\text{new}})$	predicted risk of event for X_{new}
$R_t(D_{b,m}^*)(X_{\text{oob}})$	predicted risk for subject oob

Estimate of Brier Score in right censored survival data

Inverse Probability of Censoring Weighted estimate

$$(1_{\{T_{new} \leq t\}} - R_t(X_{new}))^2 W_t(X_{new}, \hat{G})$$

where \hat{G} is an estimate of the conditional censoring distribution $G(t|X) = P(C > t|X)$ as obtained with the reverse Kaplan-Meier, an undersmoothed HAL, or a super learner for the censoring survival distribution.

The weights are given by:

$$W_t(X_{new}, \hat{G}) = \frac{1_{\{T_{new} \leq t\}} \Delta_{new}}{\hat{G}(T_{new} - |X_{new})} + \frac{1_{\{T_{new} > t\}}}{\hat{G}(t|X_{new})}$$

Leave-one-out-bootstrap ³

Algorithm:

1. Draw a bootstrap sample $D_{b,m}^*$
2. Train all the prediction models: $R_t^1(D_{b,m}^*), R_t^2(D_{b,m}^*), \dots$
3. For all out of bag subjects $i \notin D_{b,m}^*$ calculate residuals $(1_{\{T_i \leq t\}} - R_t(D_{b,m}^*)(X_i))^2$ for each prediction model
4. Repeat B times

³Efron & Tibshirani (1997)

Leave-one-out-bootstrap ³

Algorithm:

1. Draw a bootstrap sample $D_{b,m}^*$
2. Train all the prediction models: $R_t^1(D_{b,m}^*), R_t^2(D_{b,m}^*), \dots$
3. For all out of bag subjects $i \notin D_{b,m}^*$ calculate residuals $(1_{\{T_i \leq t\}} - R_t(D_{b,m}^*)(X_i))^2$ for each prediction model
4. Repeat B times

Leave-one-out bootstrap estimate

$$\hat{\mu}_{t,n} = \frac{1}{n} \sum_{i=1}^n \frac{1}{\eta_{i,B}} \sum_{b=1}^B (1_{\{T_i \leq t\}} - R_t(D_{b,m}^*)(X_i))^2 1_{\{X_i \notin D_{b,m}^*\}} W_t(X_i, \hat{G})$$

where $\eta_{i,B}$ = number of bootstrap samples in which subject i is not included.

³Efron & Tibshirani (1997)

- Q5: In a survival setting where IPCW is used to estimate performance/loss: estimate censoring distribution in each training set separately?
- A5: No, the formula for the leave-one-out bootstrap estimator clearly indicates to estimate the censoring distribution with all data.

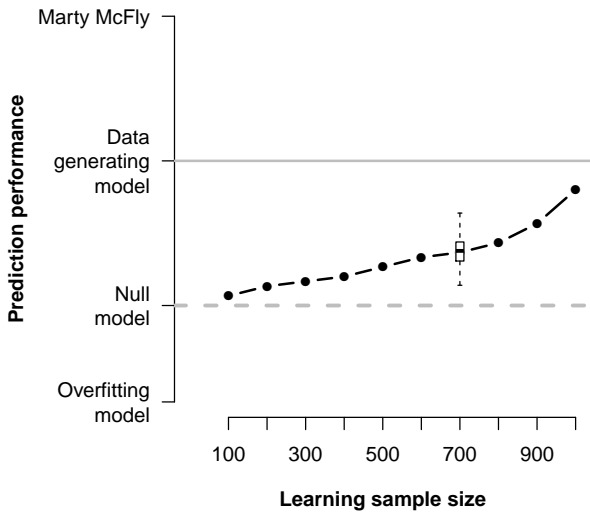
The estimated parameter

The bootstrap cross-validated estimate is an estimate of the expected value of the average residual:

$$\mathbb{E}_{new} \left[\mathbb{E}_{D_m} \left[\left(1_{\{\tilde{T}_{new} \leq t\}} - R_t(D_m)(X_{new}) \right)^2 \mid (\tilde{T}_{new}, \Delta_{new}, X_{new}) \right] \right]$$

- we are really interested in the performance of $R_t(D_n)$
- we get an estimate of the performance of the average model across all possible training data sets of size $m < n$:

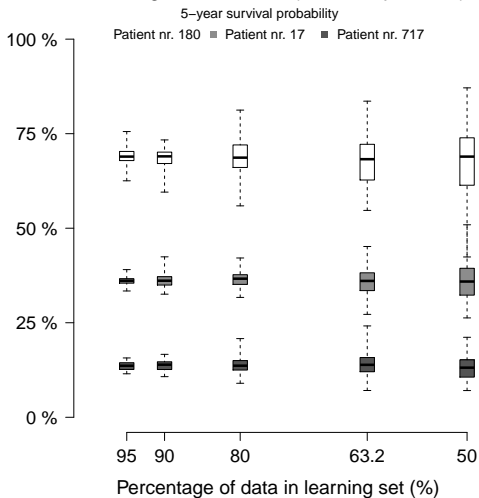
this is a bit unfortunate



Q1: 5-fold, 10-fold, 20-fold or bootstrap?

A1: The target parameter depends on this choice. There is a bias-variance tradeoff ...

Cox regression model (B=100 repetitions)



Discrete super learner

```
ff <- Surv(time,event)~ age + sex+ bili + protime +
  albumin + chol +copper +ast +platelet+spiders+trig+
  hepato+ascites
# Cox
fit.cox <- list(
  lasso=penalizedS3(ff,data=pbcc, model="cox", trace=0,
    lambda2=0, lambda1=4),
  elastic=penalizedS3(ff,data=pbcc, model="cox", trace=0,
    lambda2=5, lambda1=18))
# random forest
fit.forest=list(
  Forest5=rfsrc(ff,data=pbcc, seed=8, nodesize=5, ntree=100,
    mtry=13),
  Forest43=rfsrc(ff,data=pbcc, seed=8, nodesize=43, ntree=100)
)
# leave-one-out estimate of performance
set.seed(13)
x=Score(c(fit.cox, fit.forest),
  data=pbcc, B=100, M=.632*NROW(pbcc),
  split.method="loob",
  formula=Surv(time, status!=1)~1,
  times=1000, summary="risk")
summary(x, what="score")
```

Discrete super learner

times	Model	AUC (%)	Brier (%)
1000	Null model	50.0	15.3 [12.4;18.3]
1000	Null model	50.0 [50.0;50.0]	15.3 [12.4;18.3]
1000	lasso	86.0 [80.7;91.4]	11.2 [8.5;13.8]
1000	elastic	84.7 [79.2;90.1]	11.6 [8.9;14.3]
1000	Forest5	84.4 [79.0;89.8]	11.8 [9.4;14.2]
1000	Forest43	86.4 [81.3;91.6]	11.6 [9.4;13.8]

The winner is the lasso.

But,

- the result is quite sensitive to the seed
- the tuning parameters were not optimized in the training phase within the folds (to save computation time).

Ensemble super learner: level-1 data

ID	time	status	Null model	lasso	elastic	Forest5	Forest43
234	3255	1	17.9	2.7	4.1	4.1	6.5
169	2241	0	18.5	3.0	3.9	0.8	6.7
75	1230	1	18.4	3.2	4.4	0.4	6.6
158	2055	1	19.0	3.2	5.6	0.1	7.0
188	2449	1	18.2	3.2	6.6	23.2	8.0
11	198	1	18.0	3.4	6.1	0.4	7.5
261	4039	1	19.3	3.4	5.6	5.1	8.1
79	1271	1	18.3	3.5	5.3	0.2	6.7
128	1701	1	18.5	3.5	5.2	0.4	6.7
72	1212	1	18.3	3.6	6.7	7.2	8.1
163	2176	1	18.3	3.6	6.1	0.3	6.8
103	1434	1	18.4	3.8	4.7	0.4	6.7
...							
...							

Ensemble super learner

Breiman (1996): *Wolpert's idea is that the level one data has more information in it, and can be used to construct "good" combinations of learners.*

But he also remarks that just how to use level one data to form accurate combinations is "black art".

Breiman proposed ridge regression

```
penalized(Surv(time,status)~NullModel+lasso+elastic+
  forest5+forest43,
  data=level1,
  lambda2=7)
```

Q6: Use prediction performance to create a new learner?

A6: Yes, but ...

- ... the ensemble super learner has a tuning parameter (ridge L_2 penalty)
- ... an outer cross-validation loop is required to assess the performance of the super learner.
- ... it is possible to create a hyper learner by combining super learners and so on and so forth
- ... garbage in garbage out